

```
! pip install ptfllops
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting ptfllops
  Downloading ptfllops-0.6.9.tar.gz (12 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: torch in /usr/local/lib/python3.8/dist-packages (from ptfllops) (1.13.1+cull6)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torch->ptfllops) (4.5.0)
Building wheels for collected packages: ptfllops
  Building wheel for ptfllops (setup.py) ... done
  Created wheel for ptfllops: filename=ptfllops-0.6.9-py3-none-any.whl size=11712 sha256=573aa52379379f15b60def1503efad9c768691422b979d93eb9cc6745516a053
  Stored in directory: /root/.cache/pip/wheels/b6/86/d5/cf62a3571b005f91cd9accefc5e10f40214538be997198afad
Successfully built ptfllops
Installing collected packages: ptfllops
Successfully installed ptfllops-0.6.9
```

```
import torch
import torchvision
from torch import nn
from torchvision import transforms
import torch.optim as optim
from torchsummary import summary
from ptfllops import get_model_complexity_info
import matplotlib.pyplot as plt
import numpy as np
```

```
# Define how we want images transformed
resize = (28, 28)
trans = transforms.Compose([transforms.Resize(resize),
                           transforms.ToTensor()])
```

```
# Create training and validation sets
training_set = torchvision.datasets.FashionMNIST('./data', train=True,
                                                transform=trans, download=True)
validation_set = torchvision.datasets.FashionMNIST('./data', train=False,
                                                  transform=trans, download=True)
```

```
# Create dataloaders for each set
training_loader = torch.utils.data.DataLoader(training_set, batch_size=32,
                                              shuffle=True, num_workers=2)
validation_loader = torch.utils.data.DataLoader(validation_set, batch_size=32,
                                              shuffle=False, num_workers=2)
```

```
classes = ('T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
           'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle Boot')
```

```
print("Training set size:", len(training_set))
print("Validation set size:", len(validation_set))
```

```

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz to ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz
100% 26421880/26421880 [00:01<00:00, 26287683.24it/s]
Extracting ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz to ./data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw/train-labels-idx1-ubyte.gz
100% 29515/29515 [00:00<00:00, 298289.42it/s]
Extracting ./data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to ./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz
100% 4422102/4422102 [00:00<00:00, 7200833.18it/s]
Extracting ./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to ./data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz
100% 5148/5148 [00:00<00:00, 84372.85it/s]
Extracting ./data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/FashionMNIST/raw

```

▼ 1) MaxNet w/ ReLU

```

MaxLeNet = nn.Sequential(

    nn.Conv2d(1, 6, kernel_size=5, padding=2), nn.ReLU(),          #LazyConv2d(6, kernel_size=5, padding=2), nn.Sigmoid(),
    nn.MaxPool2d(2),                                              #nn.AvgPool2d(kernel_size=2, stride=2),
    nn.Conv2d(6, 16, kernel_size=5), nn.ReLU(),                  #LazyConv2d(16, kernel_size=5), nn.Sigmoid(),
    nn.MaxPool2d(2),                                              #nn.AvgPool2d(kernel_size=2, stride=2),
    nn.Flatten(),
    nn.Linear(400, 120), nn.ReLU(),                                #LazyLinear(120), nn.Sigmoid(),
    nn.Linear(120, 84), nn.ReLU(),                                #LazyLinear(84), nn.Sigmoid(),
    nn.Linear(84, 10)                                             #LazyLinear(num_classes))

)

summary(MaxLeNet, input_size = (1, 28, 28), batch_size = 32)

```

Layer (type)	Output Shape	Param #
-----	-----	-----
Conv2d-1	[32, 6, 28, 28]	156
ReLU-2	[32, 6, 28, 28]	0
MaxPool2d-3	[32, 6, 14, 14]	0
Conv2d-4	[32, 16, 10, 10]	2,416
ReLU-5	[32, 16, 10, 10]	0
MaxPool2d-6	[32, 16, 5, 5]	0
Flatten-7	[32, 400]	0
Linear-8	[32, 120]	48,120
ReLU-9	[32, 120]	0
Linear-10	[32, 84]	10,164
ReLU-11	[32, 84]	0
Linear-12	[32, 10]	850
=====	=====	=====

```

Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0
-----
Input size (MB): 0.10
Forward/backward pass size (MB): 3.66
Params size (MB): 0.24
Estimated Total Size (MB): 3.99
-----

```

```
# --- APPLY TO ALL EXPERIMENTS ---
```

```

macs, params = get_model_complexity_info(MaxLeNet, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)

```

```
print('Computational complexity: ', macs)
```

```
print('Number of parameters: ', params)
```

```
Warning: module Flatten is treated as a zero-op.
```

```

Sequential(
  61.71 k, 100.000% Params, 435.85 KMac, 100.000% MACs,
  (0): Conv2d(156, 0.253% Params, 122.3 KMac, 28.061% MACs, 1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 4.7 KMac, 1.079% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 4.7 KMac, 1.079% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(2.42 k, 3.915% Params, 241.6 KMac, 55.432% MACs, 6, 16, kernel_size=(5, 5), stride=(1, 1))
  (4): ReLU(0, 0.000% Params, 1.6 KMac, 0.367% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 1.6 KMac, 0.367% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (7): Linear(48.12 k, 77.983% Params, 48.12 KMac, 11.040% MACs, in_features=400, out_features=120, bias=True)
  (8): ReLU(0, 0.000% Params, 120.0 Mac, 0.028% MACs, )
  (9): Linear(10.16 k, 16.472% Params, 10.16 KMac, 2.332% MACs, in_features=120, out_features=84, bias=True)
  (10): ReLU(0, 0.000% Params, 84.0 Mac, 0.019% MACs, )
  (11): Linear(850, 1.377% Params, 850.0 Mac, 0.195% MACs, in_features=84, out_features=10, bias=True)
)
Computational complexity: 435.85 KMac
Number of parameters: 61.71 k

```

```
# Define the training loop for each epoch
```

```
def trainLoop(dataloader, model, loss_fn, optimizer):
```

```

    numBatches = len(dataloader)
    dataSize = len(dataloader.dataset)
    totalLoss = 0
    numCorrect = 0

```

```
    for batch, (X, y) in enumerate(dataloader):
```

```

        pred = model(X)
        loss = loss_fn(pred, y)

```

```

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

```

```
        totalLoss = totalLoss + loss.item()
```

```

    if batch % 1000 == 0:
        loss = loss.item()

```

```

        interLosses.append(loss)
        avgLoss = totalLoss / (batch + 1)
        avgLosses.append(avgLoss)
        print("loss:", loss)

    pred = model(X)
    numCorrect = numCorrect + (pred.argmax(1) == y).type(torch.float).sum().item()

trainAcc = numCorrect / dataSize
trainHist.append(trainAcc)

epochLoss = totalLoss / len(dataloader)
trainLosses.append(epochLoss)

# Define the validation loop for each epoch
def valLoop(dataloader, model, loss_fn):

    numBatches = len(dataloader)
    dataSize = len(dataloader.dataset)
    valLoss = 0
    numCorrect = 0

    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            valLoss = valLoss + loss_fn(pred, y).item()
            numCorrect = numCorrect + (pred.argmax(1) == y).type(torch.float).sum().item()

    valAcc = numCorrect / dataSize
    valHist.append(valAcc)
    valAccPercent = valAcc * 100

    avgLoss = valLoss / numBatches
    valLosses.append(avgLoss)

    print("Validation Accuracy:", valAccPercent, "    Validation Loss: ", avgLoss)
    print(" ")

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(MaxLeNet.parameters(), lr=0.001)

# Begin training over 20 epochs
epochs = 20
valHist = []
valLosses = []
trainHist = []
trainLosses = []
interLosses = []
avgLosses = []

for t in range(epochs):
    actualEpoch = t+1

```

```
print("Epoch", actualEpoch)
trainLoop(training_loader, MaxLeNet, criterion, optimizer)
valLoop(validation_loader, MaxLeNet, criterion)
-----
Validation Accuracy: 89.74      Validation Loss: 0.29269843495572906

Epoch 10
loss: 0.26472964882850647
loss: 0.058733873069286346
Validation Accuracy: 90.14999999999999      Validation Loss: 0.27200968651630625

Epoch 11
loss: 0.059090666472911835
loss: 0.15066036581993103
Validation Accuracy: 90.22      Validation Loss: 0.26984099586741234

Epoch 12
loss: 0.175387904047966
loss: 0.14507068693637848
Validation Accuracy: 90.46      Validation Loss: 0.27179381690514737

Epoch 13
loss: 0.17627651989459991
loss: 0.21958458423614502
Validation Accuracy: 90.55      Validation Loss: 0.27978623022857946

Epoch 14
loss: 0.13553816080093384
loss: 0.17807172238826752
Validation Accuracy: 90.07      Validation Loss: 0.2844237531550205

Epoch 15
loss: 0.07433485239744186
loss: 0.03841002658009529
Validation Accuracy: 90.62      Validation Loss: 0.28944059723791793

Epoch 16
loss: 0.025794582441449165
loss: 0.03645904362201691
Validation Accuracy: 90.38000000000001      Validation Loss: 0.2970541306494619

Epoch 17
loss: 0.09444670379161835
loss: 0.2125445306301117
Validation Accuracy: 90.28      Validation Loss: 0.3068658906442765

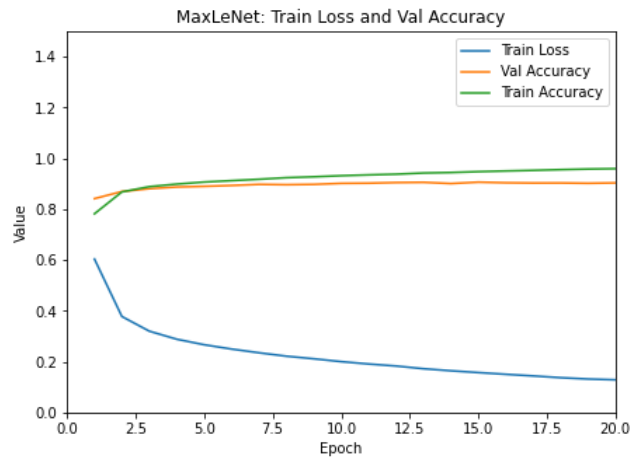
Epoch 18
loss: 0.3425464332103729
loss: 0.22572341561317444
Validation Accuracy: 90.3      Validation Loss: 0.3051984432799081

Epoch 19
loss: 0.10637159645557404
loss: 0.1293027549982071
Validation Accuracy: 90.18      Validation Loss: 0.31781230903804875

Epoch 20
loss: 0.18390409648418427
loss: 0.14328217506408691
Validation Accuracy: 90.33      Validation Loss: 0.31958482676325517
```

```
#trainLosses, valHist, interLosses, avgLosses
```

```
# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("MaxLeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()
```



2a) Experimental LeNet - Window/Kernel Size

```
ExpLeNetWind = nn.Sequential(
    nn.Conv2d(1, 6, kernel_size=3, padding=2), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Conv2d(6, 16, kernel_size=3), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Flatten(),
    nn.Linear(576, 120), nn.ReLU(),
    nn.Linear(120, 84), nn.ReLU(),
    nn.Linear(84, 10)
)
```

```
summary(ExpLeNetWind, input_size = (1, 28, 28), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 6, 30, 30]	60
ReLU-2	[32, 6, 30, 30]	0
MaxPool2d-3	[32, 6, 15, 15]	0
Conv2d-4	[32, 16, 13, 13]	880
ReLU-5	[32, 16, 13, 13]	0
MaxPool2d-6	[32, 16, 6, 6]	0
Flatten-7	[32, 576]	0
Linear-8	[32, 120]	69,240
ReLU-9	[32, 120]	0
Linear-10	[32, 84]	10,164
ReLU-11	[32, 84]	0
Linear-12	[32, 10]	850
Total params: 81,194		
Trainable params: 81,194		
Non-trainable params: 0		
Input size (MB): 0.10		
Forward/backward pass size (MB): 4.67		
Params size (MB): 0.31		
Estimated Total Size (MB): 5.08		

```
macs, params = get_model_complexity_info(ExpLeNetWind, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)
print('Computational complexity: ', macs)
print('Number of parameters: ', params)
```

```
Warning: module Flatten is treated as a zero-op.
Sequential(
```

```
  81.19 k, 100.000% Params, 299.39 KMac, 100.000% MACs,
  (0): Conv2d(60, 0.074% Params, 54.0 KMac, 18.037% MACs, 1, 6, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 5.4 KMac, 1.804% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 5.4 KMac, 1.804% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(880, 1.084% Params, 148.72 KMac, 49.675% MACs, 6, 16, kernel_size=(3, 3), stride=(1, 1))
  (4): ReLU(0, 0.000% Params, 2.7 KMac, 0.903% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 2.7 KMac, 0.903% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (7): Linear(69.24 k, 85.277% Params, 69.24 KMac, 23.127% MACs, in_features=576, out_features=120, bias=True)
  (8): ReLU(0, 0.000% Params, 120.0 Mac, 0.040% MACs, )
  (9): Linear(10.16 k, 12.518% Params, 10.16 KMac, 3.395% MACs, in_features=120, out_features=84, bias=True)
  (10): ReLU(0, 0.000% Params, 84.0 Mac, 0.028% MACs, )
  (11): Linear(850, 1.047% Params, 850.0 Mac, 0.284% MACs, in_features=84, out_features=10, bias=True)
)
```

```
Computational complexity: 299.39 KMac
Number of parameters: 81.19 k
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(ExpLeNetWind.parameters(), lr=0.001)
```

```

# Begin training over 20 epochs
epochs = 20
valHist = []
valLosses = []
trainHist = []
trainLosses = []
interLosses = []
avgLosses = []

for t in range(epochs):
    actualEpoch = t+1
    print("Epoch", actualEpoch)
    trainLoop(training_loader, ExpLeNetWind, criterion, optimizer)
    valLoop(validation_loader, ExpLeNetWind, criterion)
    #####
    Validation Accuracy: 88.85      Validation Loss:  0.2982958235358373

    Epoch 10
    loss: 0.28388839960098267
    loss: 0.14293548464775085
    Validation Accuracy: 90.16999999999999      Validation Loss:  0.267375189001663

    Epoch 11
    loss: 0.1129353865981102
    loss: 0.26229333877563477
    Validation Accuracy: 90.33      Validation Loss:  0.264691455861012

    Epoch 12
    loss: 0.14219126105308533
    loss: 0.2240721434354782
    Validation Accuracy: 89.97      Validation Loss:  0.28073896977086416

    Epoch 13
    loss: 0.2418632209300995
    loss: 0.17405880987644196
    Validation Accuracy: 90.51      Validation Loss:  0.2779467262970373

    Epoch 14
    loss: 0.11901608854532242
    loss: 0.04241831600666046
    Validation Accuracy: 90.16      Validation Loss:  0.2886403740678256

    Epoch 15
    loss: 0.21292787790298462
    loss: 0.08505263179540634
    Validation Accuracy: 90.4      Validation Loss:  0.27278012166412685

    Epoch 16
    loss: 0.0945063978433609
    loss: 0.24951516091823578
    Validation Accuracy: 90.47      Validation Loss:  0.2857096066799598

    Epoch 17
    loss: 0.18068593740463257
    loss: 0.049970388412475586
    Validation Accuracy: 90.49000000000001      Validation Loss:  0.2936227724325067

    Epoch 18
    loss: 0.18012750148773193

```



```

loss: 0.180588960647583
Validation Accuracy: 90.66      Validation Loss: 0.28525481897349747

Epoch 19
loss: 0.13082930445671082
loss: 0.22957293689250946
Validation Accuracy: 90.71000000000001      Validation Loss: 0.2929377834291789

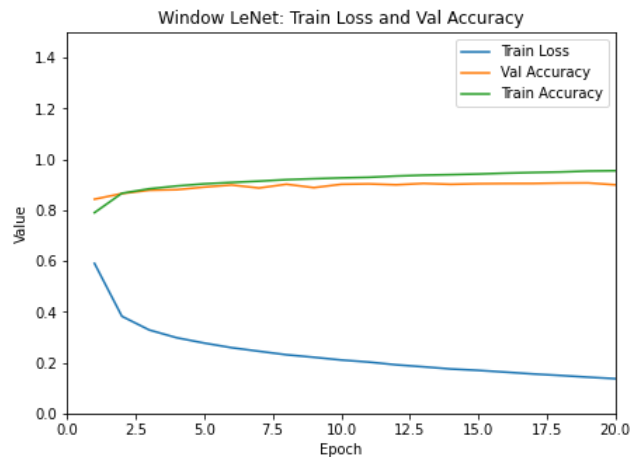
Epoch 20
loss: 0.10809242725372314
loss: 0.21371260285377502
Validation Accuracy: 89.96      Validation Loss: 0.3274085383398083

```

```

# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Window LeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()

```



2b) Experimental LeNet - Output Channels

```

ExpLeNetChan = nn.Sequential(
    nn.Conv2d(1, 32, kernel_size=5, padding=2), nn.ReLU(),

```

```

        nn.MaxPool2d(2),
        nn.Conv2d(32, 32, kernel_size=5), nn.ReLU(),
        nn.MaxPool2d(2),
        nn.Flatten(),
        nn.Linear(800, 256), nn.ReLU(),
        nn.Linear(256, 64), nn.ReLU(),
        nn.Linear(64, 10)
    )

summary(ExpLeNetChan, input_size = (1, 28, 28), batch_size =32)

```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 32, 28, 28]	832
ReLU-2	[32, 32, 28, 28]	0
MaxPool2d-3	[32, 32, 14, 14]	0
Conv2d-4	[32, 32, 10, 10]	25,632
ReLU-5	[32, 32, 10, 10]	0
MaxPool2d-6	[32, 32, 5, 5]	0
Flatten-7	[32, 800]	0
Linear-8	[32, 256]	205,056
ReLU-9	[32, 256]	0
Linear-10	[32, 64]	16,448
ReLU-11	[32, 64]	0
Linear-12	[32, 10]	650
Total params: 248,618		
Trainable params: 248,618		
Non-trainable params: 0		
Input size (MB): 0.10		
Forward/backward pass size (MB): 15.89		
Params size (MB): 0.95		
Estimated Total Size (MB): 16.94		

```

macs, params = get_model_complexity_info(ExpLeNetChan, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)

print('Computational complexity: ', macs)
print('Number of parameters: ', params)

```

Warning: module Flatten is treated as a zero-op.

```

Sequential(
  248.62 k, 100.000% Params, 3.49 MMac, 100.000% MACs,
  (0): Conv2d(832, 0.335% Params, 652.29 KMac, 18.666% MACs, 1, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 25.09 KMac, 0.718% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 25.09 KMac, 0.718% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(25.63 k, 10.310% Params, 2.56 MMac, 73.349% MACs, 32, 32, kernel_size=(5, 5), stride=(1, 1))
  (4): ReLU(0, 0.000% Params, 3.2 KMac, 0.092% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 3.2 KMac, 0.092% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (7): Linear(205.06 k, 82.478% Params, 205.06 KMac, 5.868% MACs, in_features=800, out_features=256, bias=True)
  (8): ReLU(0, 0.000% Params, 256.0 Mac, 0.007% MACs, )
  (9): Linear(16.45 k, 6.616% Params, 16.45 KMac, 0.471% MACs, in_features=256, out_features=64, bias=True)
  (10): ReLU(0, 0.000% Params, 64.0 Mac, 0.002% MACs, )
  (11): Linear(650, 0.261% Params, 650.0 Mac, 0.019% MACs, in_features=64, out_features=10, bias=True)
)

```

Computational complexity: 3.49 MMac
Number of parameters: 248.62 k

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(ExpLeNetChan.parameters(), lr=0.001)
```

```
# Begin training over 20 epochs  
epochs = 20  
valHist = []  
valLosses = []  
trainHist = []  
trainLosses = []  
interLosses = []  
avgLosses = []
```

```
for t in range(epochs):  
    actualEpoch = t+1  
    print("Epoch", actualEpoch)  
    trainLoop(training_loader, ExpLeNetChan, criterion, optimizer)  
    valLoop(validation_loader, ExpLeNetChan, criterion)  
  
    loss: 0.10545231401920319  
    Validation Accuracy: 91.77      Validation Loss: 0.2635595384544839  
  
    Epoch 10  
    loss: 0.13737992942333221  
    loss: 0.10598868876695633  
    Validation Accuracy: 91.14999999999999      Validation Loss: 0.2752029822704891  
  
    Epoch 11  
    loss: 0.05570889264345169  
    loss: 0.01653691940009594  
    Validation Accuracy: 91.34      Validation Loss: 0.3112751868026801  
  
    Epoch 12  
    loss: 0.0803329274058342  
    loss: 0.12001214921474457  
    Validation Accuracy: 91.23      Validation Loss: 0.30967554425338684  
  
    Epoch 13  
    loss: 0.06855528056621552  
    loss: 0.03943931683897972  
    Validation Accuracy: 91.73      Validation Loss: 0.3185548261760142  
  
    Epoch 14  
    loss: 0.1979738175868988  
    loss: 0.037628620862960815  
    Validation Accuracy: 91.47      Validation Loss: 0.325786305490298  
  
    Epoch 15  
    loss: 0.00550600653514266  
    loss: 0.02556084655225277  
    Validation Accuracy: 90.53      Validation Loss: 0.392708921642489  
  
    Epoch 16  
    loss: 0.0077915857546031475  
    loss: 0.009563696570694447
```

Validation Accuracy: 91.58 Validation Loss: 0.39452129550510273


Epoch 17
loss: 0.019824210554361343
loss: 0.11915914714336395
Validation Accuracy: 91.45 Validation Loss: 0.3854067757504043

Epoch 18
loss: 0.07205665856599808
loss: 0.03600491210818291
Validation Accuracy: 90.67 Validation Loss: 0.4403740083278059

Epoch 19
loss: 0.004822830203920603
loss: 0.01705688238143921
Validation Accuracy: 91.18 Validation Loss: 0.38978407756029704

Epoch 20
loss: 0.015321574173867702
loss: 0.007182026281952858
Validation Accuracy: 91.18 Validation Loss: 0.47341194060242797

```
# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Channel LeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()
```



2c) Experimental LeNet - Conv Layers

```
ExpLeNetConv = nn.Sequential(  
  
    nn.Conv2d(1, 6, kernel_size=5, padding=2), nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Conv2d(6, 16, kernel_size=5, padding=2), nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Conv2d(16, 32, kernel_size=5, padding=2), nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Conv2d(32, 64, kernel_size=5, padding=2), nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Flatten(),  
    nn.Linear(64, 32), nn.ReLU(),  
    nn.Linear(32, 16), nn.ReLU(),  
    nn.Linear(16, 10)  
)
```

```
summary(ExpLeNetConv, input_size = (1, 28, 28), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 6, 28, 28]	156
ReLU-2	[32, 6, 28, 28]	0
MaxPool2d-3	[32, 6, 14, 14]	0
Conv2d-4	[32, 16, 14, 14]	2,416
ReLU-5	[32, 16, 14, 14]	0
MaxPool2d-6	[32, 16, 7, 7]	0
Conv2d-7	[32, 32, 7, 7]	12,832
ReLU-8	[32, 32, 7, 7]	0
MaxPool2d-9	[32, 32, 3, 3]	0
Conv2d-10	[32, 64, 3, 3]	51,264
ReLU-11	[32, 64, 3, 3]	0
MaxPool2d-12	[32, 64, 1, 1]	0
Flatten-13	[32, 64]	0
Linear-14	[32, 32]	2,080
ReLU-15	[32, 32]	0
Linear-16	[32, 16]	528
ReLU-17	[32, 16]	0
Linear-18	[32, 10]	170

=====
Total params: 69,446
Trainable params: 69,446
Non-trainable params: 0
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 5.48
Params size (MB): 0.26
Estimated Total Size (MB): 5.84
=====

```
macs, params = get_model_complexity_info(ExpLeNetConv, (1, 28, 28), as_strings=True,
```

```

        print_per_layer_stat=True, verbose=True)

print('Computational complexity: ', macs)
print('Number of parameters: ', params)

Warning: module Flatten is treated as a zero-op.
Sequential(
  69.45 k, 100.000% Params, 1.71 MMac, 100.000% MACs,
  (0): Conv2d(156, 0.225% Params, 122.3 KMac, 7.157% MACs, 1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 4.7 KMac, 0.275% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 4.7 KMac, 0.275% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(2.42 k, 3.479% Params, 473.54 KMac, 27.712% MACs, 6, 16, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (4): ReLU(0, 0.000% Params, 3.14 KMac, 0.184% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 3.14 KMac, 0.184% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(12.83 k, 18.478% Params, 628.77 KMac, 36.796% MACs, 16, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (7): ReLU(0, 0.000% Params, 1.57 KMac, 0.092% MACs, )
  (8): MaxPool2d(0, 0.000% Params, 1.57 KMac, 0.092% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (9): Conv2d(51.26 k, 73.819% Params, 461.38 KMac, 27.000% MACs, 32, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (10): ReLU(0, 0.000% Params, 576.0 Mac, 0.034% MACs, )
  (11): MaxPool2d(0, 0.000% Params, 576.0 Mac, 0.034% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (12): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (13): Linear(2.08 k, 2.995% Params, 2.08 KMac, 0.122% MACs, in_features=64, out_features=32, bias=True)
  (14): ReLU(0, 0.000% Params, 32.0 Mac, 0.002% MACs, )
  (15): Linear(528, 0.760% Params, 528.0 Mac, 0.031% MACs, in_features=32, out_features=16, bias=True)
  (16): ReLU(0, 0.000% Params, 16.0 Mac, 0.001% MACs, )
  (17): Linear(170, 0.245% Params, 170.0 Mac, 0.010% MACs, in_features=16, out_features=10, bias=True)
)
Computational complexity:  1.71 MMac
Number of parameters:  69.45 k

```

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(ExpLeNetConv.parameters(), lr=0.001)

```

```

# Begin training over 20 epochs

```

```

epochs = 20
valHist = []
valLosses = []
trainHist = []
trainLosses = []
interLosses = []
avgLosses = []

```

```

for t in range(epochs):
    actualEpoch = t+1
    print("Epoch", actualEpoch)
    trainLoop(training_loader, ExpLeNetConv, criterion, optimizer)
    valLoop(validation_loader, ExpLeNetConv, criterion)

    loss: 0.18912243843078613
    Validation Accuracy: 89.59      Validation Loss: 0.30505555555129205

    Epoch 10
    loss: 0.22251185774803162
    loss: 0.1909756362438202
    Validation Accuracy: 89.67      Validation Loss: 0.29852879498475277

    Epoch 11
    loss: 0.05390220135450363
    -

```

```

loss: 0.5136352181434631
Validation Accuracy: 89.86      Validation Loss: 0.2922220514795651

Epoch 12
loss: 0.30488497018814087
loss: 0.16639462113380432
Validation Accuracy: 90.01      Validation Loss: 0.2882175712992018

Epoch 13
loss: 0.07017980515956879
loss: 0.16630296409130096
Validation Accuracy: 89.42999999999999      Validation Loss: 0.3089518357818119

Epoch 14
loss: 0.21424993872642517
loss: 0.05710185319185257
Validation Accuracy: 90.12      Validation Loss: 0.303217229138786

Epoch 15
loss: 0.512250542640686
loss: 0.2508762776851654
Validation Accuracy: 90.29      Validation Loss: 0.3028508698026212

Epoch 16
loss: 0.16260407865047455
loss: 0.22891202569007874
Validation Accuracy: 89.96      Validation Loss: 0.3137348194639332

Epoch 17
loss: 0.05506058782339096
loss: 0.1931689828634262
Validation Accuracy: 90.28      Validation Loss: 0.3300599384660157

Epoch 18
loss: 0.09002522379159927
loss: 0.23958276212215424
Validation Accuracy: 90.16      Validation Loss: 0.33050147842127864

Epoch 19
loss: 0.046834371984004974
loss: 0.12266209721565247
Validation Accuracy: 89.92      Validation Loss: 0.3062375673827843

Epoch 20
loss: 0.07435212284326553
loss: 0.08859160542488098
Validation Accuracy: 90.02      Validation Loss: 0.3225944591656375


```

```

# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Conv Layer LeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()

```

```
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()
```



2d) Experimental LeNet - FC Layers

```
ExpLeNetFC = nn.Sequential(
    nn.Conv2d(1, 6, kernel_size=5, padding=2), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Conv2d(6, 16, kernel_size=5), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Flatten(),
    nn.Linear(400, 256), nn.ReLU(),
    nn.Linear(256, 128), nn.ReLU(),
    nn.Linear(128, 64), nn.ReLU(),
    nn.Linear(64, 32), nn.ReLU(),
    nn.Linear(32, 10)
)
```

```
summary(ExpLeNetFC, input_size = (1, 28, 28), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 6, 28, 28]	156
ReLU-2	[32, 6, 28, 28]	0
MaxPool2d-3	[32, 6, 14, 14]	0
Conv2d-4	[32, 16, 10, 10]	2,416
ReLU-5	[32, 16, 10, 10]	0
MaxPool2d-6	[32, 16, 5, 5]	0

Flatten-7	[32, 400]	0
Linear-8	[32, 256]	102,656
ReLU-9	[32, 256]	0
Linear-10	[32, 128]	32,896
ReLU-11	[32, 128]	0
Linear-12	[32, 64]	8,256
ReLU-13	[32, 64]	0
Linear-14	[32, 32]	2,080
ReLU-15	[32, 32]	0
Linear-16	[32, 10]	330

```

=====
Total params: 148,790
Trainable params: 148,790
Non-trainable params: 0
-----

```

```

Input size (MB): 0.10
Forward/backward pass size (MB): 3.80
Params size (MB): 0.57
Estimated Total Size (MB): 4.46
-----

```

```

macs, params = get_model_complexity_info(ExpLeNetFC, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)

print('Computational complexity: ', macs)
print('Number of parameters: ', params)

```

Warning: module Flatten is treated as a zero-op.

```

Sequential(
  148.79 k, 100.000% Params, 523.21 KMac, 100.000% MACs,
  (0): Conv2d(156, 0.105% Params, 122.3 KMac, 23.376% MACs, 1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 4.7 KMac, 0.899% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 4.7 KMac, 0.899% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(2.42 k, 1.624% Params, 241.6 KMac, 46.176% MACs, 6, 16, kernel_size=(5, 5), stride=(1, 1))
  (4): ReLU(0, 0.000% Params, 1.6 KMac, 0.306% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 1.6 KMac, 0.306% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (7): Linear(102.66 k, 68.994% Params, 102.66 KMac, 19.620% MACs, in_features=400, out_features=256, bias=True)
  (8): ReLU(0, 0.000% Params, 256.0 Mac, 0.049% MACs, )
  (9): Linear(32.9 k, 22.109% Params, 32.9 KMac, 6.287% MACs, in_features=256, out_features=128, bias=True)
  (10): ReLU(0, 0.000% Params, 128.0 Mac, 0.024% MACs, )
  (11): Linear(8.26 k, 5.549% Params, 8.26 KMac, 1.578% MACs, in_features=128, out_features=64, bias=True)
  (12): ReLU(0, 0.000% Params, 64.0 Mac, 0.012% MACs, )
  (13): Linear(2.08 k, 1.398% Params, 2.08 KMac, 0.398% MACs, in_features=64, out_features=32, bias=True)
  (14): ReLU(0, 0.000% Params, 32.0 Mac, 0.006% MACs, )
  (15): Linear(330, 0.222% Params, 330.0 Mac, 0.063% MACs, in_features=32, out_features=10, bias=True)
)
Computational complexity: 523.21 KMac
Number of parameters: 148.79 k

```

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(ExpLeNetFC.parameters(), lr=0.001)

```

```

# Begin training over 20 epochs
epochs = 20
valHist = []
valLosses = []
trainHist = []

```

```
trainLosses = []
interLosses = []
avgLosses = []

for t in range(epochs):
    actualEpoch = t+1
    print("Epoch", actualEpoch)
    trainLoop(training_loader, ExpLeNetFC, criterion, optimizer)
    valLoop(validation_loader, ExpLeNetFC, criterion)

    Epoch 1
    loss: 2.290706157684326
    loss: 0.307341068983078
    Validation Accuracy: 82.17999999999999    Validation Loss: 0.48632207193884985

    Epoch 2
    loss: 0.4437580108642578
    loss: 0.35284027457237244
    Validation Accuracy: 86.22999999999999    Validation Loss: 0.38711444512247656

    Epoch 3
    loss: 0.37868165969848633
    loss: 0.2912781834602356
    Validation Accuracy: 85.41    Validation Loss: 0.397343154841909

    Epoch 4
    loss: 0.4657308757305145
    loss: 0.3759935200214386
    Validation Accuracy: 87.35000000000001    Validation Loss: 0.3495084028750563

    Epoch 5
    loss: 0.17693762481212616
    loss: 0.2395053654909134
    Validation Accuracy: 88.24    Validation Loss: 0.3206745018831457

    Epoch 6
    loss: 0.28645333647727966
    loss: 0.15221655368804932
    Validation Accuracy: 88.75    Validation Loss: 0.3223715439653054

    Epoch 7
    loss: 0.5499658584594727
    loss: 0.26621702313423157
    Validation Accuracy: 89.60000000000001    Validation Loss: 0.2914217749116615

    Epoch 8
    loss: 0.1569925844669342
    loss: 0.20477882027626038
    Validation Accuracy: 89.9    Validation Loss: 0.29118004319862056


    Epoch 9
    loss: 0.3153538107872009
    loss: 0.14894960820674896
    Validation Accuracy: 89.82    Validation Loss: 0.2950234032286623

    Epoch 10
    loss: 0.22837993502616882
    loss: 0.26834988594055176
    Validation Accuracy: 89.84    Validation Loss: 0.3026259756935671
```

```
Epoch 11
loss: 0.2310730218887329
loss: 0.11507314443588257
Validation Accuracy: 90.41      Validation Loss: 0.28189363941169393
```

```
Epoch 12
loss: 0.13440990447998047
loss: 0.28459158539772034
```

```
# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("FC Layer LeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()
```



2e) Experimental LeNet - Learning Rate

```
ExpLeNetLearn = nn.Sequential(
    nn.Conv2d(1, 6, kernel_size=5, padding=2), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Conv2d(6, 16, kernel_size=5), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Flatten(),
    nn.Linear(400, 120), nn.ReLU(),
```

```

nn.Linear(120, 84), nn.ReLU(),
nn.Linear(84, 10)
)

```

```
summary(ExpLeNetLearn, input_size = (1, 28, 28), batch_size =32)
```

```

-----
Layer (type)           Output Shape          Param #
-----
Conv2d-1               [32, 6, 28, 28]      156
ReLU-2                 [32, 6, 28, 28]       0
MaxPool2d-3            [32, 6, 14, 14]       0
Conv2d-4                [32, 16, 10, 10]     2,416
ReLU-5                 [32, 16, 10, 10]      0
MaxPool2d-6            [32, 16, 5, 5]        0
Flatten-7              [32, 400]             0
Linear-8                [32, 120]             48,120
ReLU-9                 [32, 120]             0
Linear-10               [32, 84]              10,164
ReLU-11                [32, 84]              0
Linear-12               [32, 10]              850
=====
Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0

-----
Input size (MB): 0.10
Forward/backward pass size (MB): 3.66
Params size (MB): 0.24
Estimated Total Size (MB): 3.99
-----

```

```

macs, params = get_model_complexity_info(ExpLeNetLearn, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)

print('Computational complexity: ', macs)
print('Number of parameters: ', params)

```

```

Warning: module Flatten is treated as a zero-op.
Sequential(
  61.71 k, 100.000% Params, 435.85 KMac, 100.000% MACs,
  (0): Conv2d(156, 0.253% Params, 122.3 KMac, 28.061% MACs, 1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 4.7 KMac, 1.079% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 4.7 KMac, 1.079% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(2.42 k, 3.915% Params, 241.6 KMac, 55.432% MACs, 6, 16, kernel_size=(5, 5), stride=(1, 1))
  (4): ReLU(0, 0.000% Params, 1.6 KMac, 0.367% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 1.6 KMac, 0.367% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (7): Linear(48.12 k, 77.983% Params, 48.12 KMac, 11.040% MACs, in_features=400, out_features=120, bias=True)
  (8): ReLU(0, 0.000% Params, 120.0 Mac, 0.028% MACs, )
  (9): Linear(10.16 k, 16.472% Params, 10.16 KMac, 2.332% MACs, in_features=120, out_features=84, bias=True)
  (10): ReLU(0, 0.000% Params, 84.0 Mac, 0.019% MACs, )
  (11): Linear(850, 1.377% Params, 850.0 Mac, 0.195% MACs, in_features=84, out_features=10, bias=True)
)
Computational complexity: 435.85 KMac
Number of parameters: 61.71 k

```

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(ExpLeNetLearn.parameters(), lr=0.0001)

# Begin training over 20 epochs
epochs = 20
valHist = []
valLosses = []
trainHist = []
trainLosses = []
interLosses = []
avgLosses = []

for t in range(epochs):
    actualEpoch = t+1
    print("Epoch", actualEpoch)
    trainLoop(training_loader, ExpLeNetLearn, criterion, optimizer)
    valLoop(validation_loader, ExpLeNetLearn, criterion)

    loss: 0.46455809473991394
    Validation Accuracy: 85.08      Validation Loss: 0.4200473606300811

    Epoch 10
    loss: 0.20845942199230194
    loss: 0.4548656940460205
    Validation Accuracy: 85.81      Validation Loss: 0.4020056555415596

    Epoch 11
    loss: 0.28233468532562256
    loss: 0.3258947432041168
    Validation Accuracy: 85.19      Validation Loss: 0.4096351430748408

    Epoch 12
    loss: 0.25641652941703796
    loss: 0.2643296718597412
    Validation Accuracy: 86.55000000000001      Validation Loss: 0.381915259915895

    Epoch 13
    loss: 0.29789093136787415
    loss: 0.37787431478500366
    Validation Accuracy: 85.92      Validation Loss: 0.39304941247541686

    Epoch 14
    loss: 0.3465706408023834
    loss: 0.4797326624393463
    Validation Accuracy: 86.5      Validation Loss: 0.37464668042362687

    Epoch 15
    loss: 0.4175805449485779
    loss: 0.44802987575531006
    Validation Accuracy: 86.76      Validation Loss: 0.3647009155001884

    Epoch 16
    loss: 0.43517136573791504
    loss: 0.36722302436828613
    Validation Accuracy: 87.12      Validation Loss: 0.35968529384214276

    Epoch 17
    loss: 0.15792155265808105
    loss: 0.1539457142353058
    Validation Accuracy: 87.59      Validation Loss: 0.3501290382025912

```

```

Epoch 18
loss: 0.387608140707016
loss: 0.3448188602924347
Validation Accuracy: 86.53      Validation Loss: 0.36654538631486816

Epoch 19
loss: 0.3941614329814911
loss: 0.3907373547554016
Validation Accuracy: 87.94      Validation Loss: 0.33603902344410413


Epoch 20
loss: 0.21067094802856445
loss: 0.40098559856414795
Validation Accuracy: 87.72999999999999      Validation Loss: 0.3397259250426064

```

```

# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Learn Rate LeNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()

```



```

trainLosses, trainHist

([0.9013984229405722,
 0.6214531558195749,
 0.5549947891314825,

```

```

0.5094297926028569,
0.47651880617936454,
0.451864084148407,
0.4300498764872551,
0.41432707730134327,
0.3972409782886505,
0.38560177760918934,
0.3737924686113993,
0.3616999568462372,
0.3544420238931974,
0.34659421243866284,
0.3379843934893608,
0.3310685098648071,
0.3239920284807682,
0.3173277738412221,
0.31247236766616504,
0.30721893773476283],
[0.6790833333333334,
0.76885,
0.7992166666666667,
0.81935,
0.8323833333333334,
0.8418666666666667,
0.8499,
0.8552333333333333,
0.8609666666666667,
0.8653833333333333,
0.8689333333333333,
0.8729,
0.8757166666666667,
0.8782,
0.8804666666666666,
0.8833833333333333,
0.88575,
0.8875333333333333,
0.8900333333333333,
0.8899833333333333])

```

▼ 3) Best Model w/ Dropout

```

DropLeNetChan = nn.Sequential(

    nn.Conv2d(1, 32, kernel_size=5, padding=2), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Dropout(0.3),
    nn.Conv2d(32, 32, kernel_size=5), nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Dropout(0.3),
    nn.Flatten(),
    nn.Linear(800, 256), nn.ReLU(),
    nn.Dropout(0.3),
    nn.Linear(256, 64), nn.ReLU(),
    nn.Dropout(0.3),
    nn.Linear(64, 10)

)

```

```
summary(DropLeNetChan, input_size = (1, 28, 28), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 32, 28, 28]	832
ReLU-2	[32, 32, 28, 28]	0
MaxPool2d-3	[32, 32, 14, 14]	0
Dropout-4	[32, 32, 14, 14]	0
Conv2d-5	[32, 32, 10, 10]	25,632
ReLU-6	[32, 32, 10, 10]	0
MaxPool2d-7	[32, 32, 5, 5]	0
Dropout-8	[32, 32, 5, 5]	0
Flatten-9	[32, 800]	0
Linear-10	[32, 256]	205,056
ReLU-11	[32, 256]	0
Dropout-12	[32, 256]	0
Linear-13	[32, 64]	16,448
ReLU-14	[32, 64]	0
Dropout-15	[32, 64]	0
Linear-16	[32, 10]	650
Total params: 248,618		
Trainable params: 248,618		
Non-trainable params: 0		
Input size (MB): 0.10		
Forward/backward pass size (MB): 17.70		
Params size (MB): 0.95		
Estimated Total Size (MB): 18.74		

```
macs, params = get_model_complexity_info(DropLeNetChan, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)
print('Computational complexity: ', macs)
print('Number of parameters: ', params)
```

```
Warning: module Dropout is treated as a zero-op.
Warning: module Flatten is treated as a zero-op.
Sequential(
```

```
  248.62 k, 100.000% Params, 3.49 MMac, 100.000% MACs,
  (0): Conv2d(832, 0.335% Params, 652.29 KMac, 18.666% MACs, 1, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (1): ReLU(0, 0.000% Params, 25.09 KMac, 0.718% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 25.09 KMac, 0.718% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.3, inplace=False)
  (4): Conv2d(25.63 k, 10.310% Params, 2.56 MMac, 73.349% MACs, 32, 32, kernel_size=(5, 5), stride=(1, 1))
  (5): ReLU(0, 0.000% Params, 3.2 KMac, 0.092% MACs, )
  (6): MaxPool2d(0, 0.000% Params, 3.2 KMac, 0.092% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (7): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.3, inplace=False)
  (8): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (9): Linear(205.06 k, 82.478% Params, 205.06 KMac, 5.868% MACs, in_features=800, out_features=256, bias=True)
  (10): ReLU(0, 0.000% Params, 256.0 Mac, 0.007% MACs, )
  (11): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.3, inplace=False)
  (12): Linear(16.45 k, 6.616% Params, 16.45 KMac, 0.471% MACs, in_features=256, out_features=64, bias=True)
  (13): ReLU(0, 0.000% Params, 64.0 Mac, 0.002% MACs, )
  (14): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.3, inplace=False)
  (15): Linear(650, 0.261% Params, 650.0 Mac, 0.019% MACs, in_features=64, out_features=10, bias=True)
```



```
)
Computational complexity: 3.49 MMac
Number of parameters: 248.62 k
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(DropLeNetChan.parameters(), lr=0.001)
```

```
# Begin training over 20 epochs
epochs = 20
valHist = []
valLosses = []
trainHist = []
trainLosses = []
interLosses = []
avgLosses = []
```

```
for t in range(epochs):
    actualEpoch = t+1
    print("Epoch", actualEpoch)
    trainLoop(training_loader, DropLeNetChan, criterion, optimizer)
    valLoop(validation_loader, DropLeNetChan, criterion)

    loss: 0.2781029939651489
    Validation Accuracy: 91.16      Validation Loss: 0.28653141904991275

    Epoch 10
    loss: 0.045927003026008606
    loss: 0.07794604450464249
    Validation Accuracy: 90.89      Validation Loss: 0.2941940748117888

    Epoch 11
    loss: 0.10345877707004547
    loss: 0.11044761538505554
    Validation Accuracy: 91.17      Validation Loss: 0.2977535161234367

    Epoch 12
    loss: 0.1692483127117157
    loss: 0.38319969177246094
    Validation Accuracy: 90.51      Validation Loss: 0.3777781583550282

    Epoch 13
    loss: 0.0576220266520977
    loss: 0.02092619426548481
    Validation Accuracy: 90.42      Validation Loss: 0.3792904548974821

    Epoch 14
    loss: 0.006482385098934174
    loss: 0.07958652079105377
    Validation Accuracy: 90.74      Validation Loss: 0.3468707030841384

    Epoch 15
    loss: 0.09907671809196472
    loss: 0.2174682468175888
    Validation Accuracy: 91.34      Validation Loss: 0.3576213141067174

    Epoch 16
    loss: 0.0863342434167862
    loss: 0.029391279444098473
```

```

Validation Accuracy: 91.18      Validation Loss: 0.35718619907661653

Epoch 17
loss: 0.018787117674946785
loss: 0.030273647978901863
Validation Accuracy: 90.96      Validation Loss: 0.40946935341892815

Epoch 18
loss: 0.09341385960578918
loss: 0.015946049243211746
Validation Accuracy: 91.51      Validation Loss: 0.41649855165400157

Epoch 19
loss: 0.01702222228050232
loss: 0.015503406524658203
Validation Accuracy: 91.19      Validation Loss: 0.44161497947191614

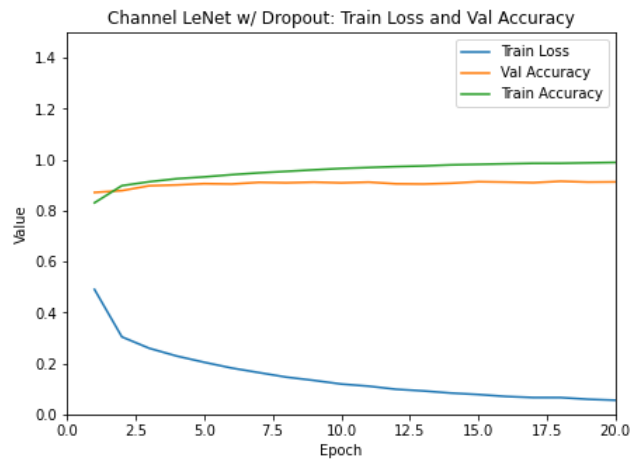
Epoch 20
loss: 0.15391159057617188
loss: 0.001051229308359325
Validation Accuracy: 91.25999999999999      Validation Loss: 0.44329784055479976

```

```

# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Channel LeNet w/ Dropout: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()

```



4) AlexNet

```
# First analyze AlexNet's complexity
AlexNet = nn.Sequential(

    nn.Conv2d(1, 96, kernel_size=11, stride=4, padding=1), nn.ReLU(),
    nn.MaxPool2d(kernel_size=3, stride=2),
    nn.Conv2d(96, 256, kernel_size=5, padding=2), nn.ReLU(),
    nn.MaxPool2d(kernel_size=3, stride=2),
    nn.Conv2d(256, 384, kernel_size=3, padding=1), nn.ReLU(),
    nn.Conv2d(384, 384, kernel_size=3, padding=1), nn.ReLU(),
    nn.Conv2d(384, 256, kernel_size=3, padding=1), nn.ReLU(),
    nn.MaxPool2d(kernel_size=3, stride=2),
    nn.Flatten(),
    nn.Linear(6400, 4096), nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(4096, 4096), nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(4096, 10)

)
```

```
summary(AlexNet, input_size = (1, 224, 224), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 96, 54, 54]	11,712
ReLU-2	[32, 96, 54, 54]	0
MaxPool2d-3	[32, 96, 26, 26]	0
Conv2d-4	[32, 256, 26, 26]	614,656
ReLU-5	[32, 256, 26, 26]	0
MaxPool2d-6	[32, 256, 12, 12]	0
Conv2d-7	[32, 384, 12, 12]	885,120
ReLU-8	[32, 384, 12, 12]	0
Conv2d-9	[32, 384, 12, 12]	1,327,488
ReLU-10	[32, 384, 12, 12]	0
Conv2d-11	[32, 256, 12, 12]	884,992
ReLU-12	[32, 256, 12, 12]	0
MaxPool2d-13	[32, 256, 5, 5]	0
Flatten-14	[32, 6400]	0
Linear-15	[32, 4096]	26,218,496
ReLU-16	[32, 4096]	0
Dropout-17	[32, 4096]	0
Linear-18	[32, 4096]	16,781,312
ReLU-19	[32, 4096]	0
Dropout-20	[32, 4096]	0
Linear-21	[32, 10]	40,970

```
=====  
Total params: 46,764,746  
Trainable params: 46,764,746  
Non-trainable params: 0  
=====
```

```
Input size (MB): 6.12
Forward/backward pass size (MB): 327.16
Params size (MB): 178.39
Estimated Total Size (MB): 511.68
-----
```

```
macs, params = get_model_complexity_info(AlexNet, (1, 224, 224), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)
```

```
print('Computational complexity: ', macs)
print('Number of parameters: ', params)
```

```
Warning: variables __flops__ or __params__ are already defined for the moduleConv2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleMaxPool2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleConv2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleMaxPool2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleConv2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleConv2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleMaxPool2d ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleLinear ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleLinear ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleReLU ptflops can affect your code!
Warning: variables __flops__ or __params__ are already defined for the moduleLinear ptflops can affect your code!
Warning: module Flatten is treated as a zero-op.
Warning: module Dropout is treated as a zero-op.
Sequential(
  46.76 M, 100.000% Params, 939.85 MMac, 100.000% MACs,
  (0): Conv2d(11.71 k, 0.025% Params, 34.15 MMac, 3.634% MACs, 1, 96, kernel_size=(11, 11), stride=(4, 4), padding=(1, 1))
  (1): ReLU(0, 0.000% Params, 279.94 KMac, 0.030% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 279.94 KMac, 0.030% MACs, kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(614.66 k, 1.314% Params, 415.51 MMac, 44.210% MACs, 96, 256, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (4): ReLU(0, 0.000% Params, 173.06 KMac, 0.018% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 173.06 KMac, 0.018% MACs, kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(885.12 k, 1.893% Params, 127.46 MMac, 13.561% MACs, 256, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (7): ReLU(0, 0.000% Params, 55.3 KMac, 0.006% MACs, )
  (8): Conv2d(1.33 M, 2.839% Params, 191.16 MMac, 20.339% MACs, 384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (9): ReLU(0, 0.000% Params, 55.3 KMac, 0.006% MACs, )
  (10): Conv2d(884.99 k, 1.892% Params, 127.44 MMac, 13.559% MACs, 384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(0, 0.000% Params, 36.86 KMac, 0.004% MACs, )
  (12): MaxPool2d(0, 0.000% Params, 36.86 KMac, 0.004% MACs, kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (13): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (14): Linear(26.22 M, 56.065% Params, 26.22 MMac, 2.790% MACs, in_features=6400, out_features=4096, bias=True)
  (15): ReLU(0, 0.000% Params, 4.1 KMac, 0.000% MACs, )
  (16): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.5, inplace=False)
  (17): Linear(16.78 M, 35.885% Params, 16.78 MMac, 1.786% MACs, in_features=4096, out_features=4096, bias=True)
  (18): ReLU(0, 0.000% Params, 4.1 KMac, 0.000% MACs, )
  (19): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.5, inplace=False)
  (20): Linear(40.97 k, 0.088% Params, 40.97 KMac, 0.004% MACs, in_features=4096, out_features=10, bias=True)
)
Computational complexity: 939.85 MMac
Number of parameters: 46.76 M
```

▼ Simplified AlexNet

```
# Simplifying AlexNet
SimpAlexNet = nn.Sequential(

    nn.Conv2d(1, 32, kernel_size=3, stride=2, padding=1), nn.ReLU(),
    nn.MaxPool2d(kernel_size=3, stride=2),
    nn.Conv2d(32, 64, kernel_size=5, padding=2), nn.ReLU(),
    nn.MaxPool2d(kernel_size=3, stride=2),
    nn.Conv2d(64, 128, kernel_size=3, padding=1), nn.ReLU(),
    nn.Conv2d(128, 128, kernel_size=3, padding=1), nn.ReLU(),
    nn.Conv2d(128, 96, kernel_size=3, padding=1), nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Flatten(),
    nn.Linear(96, 64), nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(64, 32), nn.ReLU(),
    nn.Dropout(0.5),
    nn.Linear(32, 10)

)

summary(SimpAlexNet, input_size = (1, 28, 28), batch_size =32)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[32, 32, 14, 14]	320
ReLU-2	[32, 32, 14, 14]	0
MaxPool2d-3	[32, 32, 6, 6]	0
Conv2d-4	[32, 64, 6, 6]	51,264
ReLU-5	[32, 64, 6, 6]	0
MaxPool2d-6	[32, 64, 2, 2]	0
Conv2d-7	[32, 128, 2, 2]	73,856
ReLU-8	[32, 128, 2, 2]	0
Conv2d-9	[32, 128, 2, 2]	147,584
ReLU-10	[32, 128, 2, 2]	0
Conv2d-11	[32, 96, 2, 2]	110,688
ReLU-12	[32, 96, 2, 2]	0
MaxPool2d-13	[32, 96, 1, 1]	0
Flatten-14	[32, 96]	0
Linear-15	[32, 64]	6,208
ReLU-16	[32, 64]	0
Dropout-17	[32, 64]	0
Linear-18	[32, 32]	2,080
ReLU-19	[32, 32]	0
Dropout-20	[32, 32]	0
Linear-21	[32, 10]	330

```
=====  
Total params: 392,330  
Trainable params: 392,330  
Non-trainable params: 0
```

```
-----  
Input size (MB): 0.10
```

```
Forward/backward pass size (MB): 5.34
Params size (MB): 1.50
Estimated Total Size (MB): 6.93
-----
```

```
macs, params = get_model_complexity_info(SimpAlexNet, (1, 28, 28), as_strings=True,
                                         print_per_layer_stat=True, verbose=True)
```

```
print('Computational complexity: ', macs)
```

```
print('Number of parameters: ', params)
```

```
Warning: module Flatten is treated as a zero-op.
```

```
Warning: module Dropout is treated as a zero-op.
```

```
Sequential(
  392.33 k, 100.000% Params, 3.26 MMac, 100.000% MACs,
  (0): Conv2d(320, 0.082% Params, 62.72 KMac, 1.921% MACs, 1, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (1): ReLU(0, 0.000% Params, 6.27 KMac, 0.192% MACs, )
  (2): MaxPool2d(0, 0.000% Params, 6.27 KMac, 0.192% MACs, kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(51.26 k, 13.067% Params, 1.85 MMac, 56.534% MACs, 32, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (4): ReLU(0, 0.000% Params, 2.3 KMac, 0.071% MACs, )
  (5): MaxPool2d(0, 0.000% Params, 2.3 KMac, 0.071% MACs, kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(73.86 k, 18.825% Params, 295.42 KMac, 9.050% MACs, 64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (7): ReLU(0, 0.000% Params, 512.0 Mac, 0.016% MACs, )
  (8): Conv2d(147.58 k, 37.617% Params, 590.34 KMac, 18.084% MACs, 128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (9): ReLU(0, 0.000% Params, 512.0 Mac, 0.016% MACs, )
  (10): Conv2d(110.69 k, 28.213% Params, 442.75 KMac, 13.563% MACs, 128, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(0, 0.000% Params, 384.0 Mac, 0.012% MACs, )
  (12): MaxPool2d(0, 0.000% Params, 384.0 Mac, 0.012% MACs, kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (13): Flatten(0, 0.000% Params, 0.0 Mac, 0.000% MACs, start_dim=1, end_dim=-1)
  (14): Linear(6.21 k, 1.582% Params, 6.21 KMac, 0.190% MACs, in_features=96, out_features=64, bias=True)
  (15): ReLU(0, 0.000% Params, 64.0 Mac, 0.002% MACs, )
  (16): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.5, inplace=False)
  (17): Linear(2.08 k, 0.530% Params, 2.08 KMac, 0.064% MACs, in_features=64, out_features=32, bias=True)
  (18): ReLU(0, 0.000% Params, 32.0 Mac, 0.001% MACs, )
  (19): Dropout(0, 0.000% Params, 0.0 Mac, 0.000% MACs, p=0.5, inplace=False)
  (20): Linear(330, 0.084% Params, 330.0 Mac, 0.010% MACs, in_features=32, out_features=10, bias=True)
)
Computational complexity:  3.26 MMac
Number of parameters:  392.33 k
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(SimpAlexNet.parameters(), lr=0.001)
```

```
# Begin training over 20 epochs
```

```
epochs = 20
```

```
valHist = []
```

```
valLosses = []
```

```
trainHist = []
```

```
trainLosses = []
```

```
interLosses = []
```

```
avgLosses = []
```

```
for t in range(epochs):
```

```
    actualEpoch = t+1
```

```
    print("Epoch", actualEpoch)
```

```
    trainLoop(training_loader, SimpAlexNet, criterion, optimizer)
```

```
    valLoop(validation_loader, SimpAlexNet, criterion)
```

loss: 0.18/20115/21225/39
Validation Accuracy: 88.39 Validation Loss: 0.35834061316312693

Epoch 10
loss: 0.44304436445236206
loss: 0.1485912948846817
Validation Accuracy: 88.03 Validation Loss: 0.35558546795108065

Epoch 11
loss: 0.07590296864509583
loss: 0.24297630786895752
Validation Accuracy: 89.39 Validation Loss: 0.30821049189605654

Epoch 12
loss: 0.3486815392971039
loss: 0.4212316870689392
Validation Accuracy: 88.75999999999999 Validation Loss: 0.3630248645719248

Epoch 13
loss: 0.15323634445667267
loss: 0.09813195466995239
Validation Accuracy: 88.81 Validation Loss: 0.33460332328876174

Epoch 14
loss: 0.2781636714935303
loss: 0.5736841559410095
Validation Accuracy: 89.08 Validation Loss: 0.3559204632826983

Epoch 15
loss: 0.32782769203186035
loss: 0.16226975619792938
Validation Accuracy: 89.03 Validation Loss: 0.33801869034898074

Epoch 16
loss: 0.18509599566459656
loss: 0.09188417345285416
Validation Accuracy: 89.24 Validation Loss: 0.3793918632637388


Epoch 17
loss: 0.13456179201602936
loss: 0.08670203387737274
Validation Accuracy: 88.97 Validation Loss: 0.40783253519203716

Epoch 18
loss: 0.16297142207622528
loss: 0.2289642095565796
Validation Accuracy: 89.41 Validation Loss: 0.3860520369733294

Epoch 19
loss: 0.07365482300519943
loss: 0.13822397589683533
Validation Accuracy: 88.67 Validation Loss: 0.4406382368859677

Epoch 20
loss: 0.05559298396110535
loss: 0.12641765177249908
Validation Accuracy: 89.17 Validation Loss: 0.3869175660593536

```
# Plot results
x = np.linspace(1, epochs, epochs)
plt.figure(figsize=(7,5))
plt.title("Simplified AlexNet: Train Loss and Val Accuracy")
plt.plot(x, trainLosses, label="Train Loss")
plt.plot(x, valHist, label="Val Accuracy")
plt.plot(x, trainHist, label="Train Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.legend()
plt.xlim(0, epochs)
plt.ylim(0, 1.5)
plt.show()
```



✓ 0s completed at 4:17 PM

